# Notes: IBM Bob as an AI Coding Assistant — Relevance to Project Workflow

**Status:** Reflection notes based on external training attended Aug 17, 2026 ("Build Faster with IBM Bob: AI-Assisted Coding for Hackathon Teams," Algiers). This is *not* a usage report — IBM Bob has not yet been used on any project in this portfolio. It documents what was presented, an honest comparison against the workflow actually used so far, and a concrete way to turn this from attendance into evidence.

### What was covered

IBM Bob was presented as an AI coding assistant intended to help with writing code, debugging, understanding unfamiliar frameworks, and moving from an idea to a working prototype faster, in the context of hackathon-style rapid development.

### Honest comparison to the current workflow

Every project in this portfolio (Axiom-Zero, AEGIS-Net, Vision 2030 AI Readiness Framework, ClimatePulse AI, the SEAG experiment implementation) was written and debugged without a dedicated AI coding assistant integrated into the development loop — the workflow so far has been direct implementation, testing, and manual debugging in Python.

This means there is no existing point of comparison from personal experience yet. Any claim of "IBM Bob helped me build X faster" would not be accurate, since it has not been used on any of this portfolio's code. What can be said honestly is only that the session described a category of tool (AI-assisted coding) that hasn't been part of the workflow to date.

### What this does NOT mean (explicit scope limit)

- Attending a demo of a coding assistant is not equivalent to having used one in a real project.

- No time savings, code-quality difference, or workflow improvement can be claimed until the tool is actually applied to a specific task and the outcome is compared to doing the same task without it.

#### Concrete next step (not yet done)

The most natural test case is already defined: the ClickHouse benchmark proposed in the companion notes (`clickhouse-agent-memory-eval.md`) — comparing PostgreSQL/pgvector against ClickHouse for an approval-gate audit-trail query using the SEAG experiment’s synthetic event log.

That benchmark has not been implemented yet. Using an AI coding assistant (IBM Bob or an equivalent available tool) to help implement it — and noting, honestly, whether it helped or not, and how — would turn two separate training attendances into one coherent, evidenced development cycle:

1. Attended training on ClickHouse as an agent-memory backend → documented design implications (done).
2. Attended training on AI-assisted coding → documented as a candidate tool for implementation speed (this file).
3. **Not yet done:** implement the ClickHouse benchmark, optionally using an AI coding assistant, and report what actually happened — including a negative result if the assistant didn’t meaningfully help, consistent with how negative results are reported elsewhere in this project.

Until step 3 happens, this file remains a documented consideration, not evidence of tool proficiency or workflow improvement.